


# Debezium, Kafka, and Postgres POC Documentation



## Introduction

This proof of concept (POC) demonstrates the use of **Debezium** for **Change Data Capture (CDC)** in a PostgreSQL database, capturing changes triggered by **CREATE**, **UPDATE**, or **DELETE** operations. These changes are then streamed into **Kafka** and consumed by the **baas-charges-module** (demo project) after being serialized by the **baas-analytical-module** (demo project). The consumed data is then deserialized for further processing.

This POC leverages **Docker Compose** to deploy and manage containers for each component, enabling an efficient setup and interaction between the CDC, Kafka, and data processing modules.

## Components Overview

### 1. Postgres

- Acts as the source database where data changes are captured.
- Configured with CDC (Change Data Capture) to capture row-level changes in tables.

### 2. Debezium

- An open-source CDC tool that tracks changes in databases and streams those changes to Kafka.
- Listens to the Postgres instance and streams any changes (inserts, updates, deletes) to Kafka topics.

### 3. Kafka

- Serves as the messaging backbone, where all data changes are streamed to different topics for downstream consumption.
- Kafka is configured with Zookeeper for managing its cluster metadata.

## 4. Zookeeper

- Required by Kafka to manage and coordinate distributed instances of Kafka.

## 5. Application for triggering request to the DB

- Required for CDC.

## 6. Consumer/Listener application

- Required for processing the messages from the Kafka.

# Prerequisites

Ensure **Docker** and **Docker Compose** are installed.

`docker --version`

`docker-compose --version`

# Docker Compose Configuration

The `docker-compose.yml` file below configures each component in a networked environment, enabling them to communicate with each other seamlessly.

```
version: '3.8'
services:
  zookeeper:
    image: confluentinc/cp-zookeeper:7.0.0
    container_name: zookeeper
    environment:
      ZOOKEEPER_CLIENT_PORT: 2181
      ZOOKEEPER_TICK_TIME: 2000
    ports:
      - 2181:2181

  kafka:
    image: confluentinc/cp-kafka:7.0.0
    container_name: kafka
    depends_on:
      - zookeeper
    environment:
      KAFKA_BROKER_ID: 1
      KAFKA_ZOOKEEPER_CONNECT: zookeeper:2181
      KAFKA_ADVERTISED_LISTENERS: PLAINTEXT://kafka:9092
      KAFKA_OFFSETS_TOPIC_REPLICATION_FACTOR: 1
    ports:
      - 9092:9092

  postgres:
    image: postgres:13
    container_name: postgres
```

```
environment:
  POSTGRES_USER: postgres
  POSTGRES_PASSWORD: password
  POSTGRES_DB: testdb
ports:
  - 5432:5432
volumes:
  - ./init.sql:/docker-entrypoint-initdb.d/init.sql:ro
logging:
  options:
    max-size: "10m"
    max-file: "3"

debezium:
  image: debezium/connect:1.8
  container_name: debezium
  depends_on:
    - kafka
    - postgres
  environment:
    BOOTSTRAP_SERVERS: kafka:9092
    GROUP_ID: 1
    CONFIG_STORAGE_TOPIC: debezium-connect-configs
    OFFSET_STORAGE_TOPIC: debezium-connect-offsets
    STATUS_STORAGE_TOPIC: debezium-connect-status
  ports:
    - 8083:8083
  logging:
    options:
      max-size: "10m"
      max-file: "3"
```

## Explanation of `docker-compose.yml` Configurations

### Zookeeper

- **Image:** confluentinc/cp-zookeeper:7.0.0 – This is the Zookeeper image required by Kafka.
- **Environment Variables:**
  - ZOOKEEPER\_CLIENT\_PORT: Port for Zookeeper clients to connect to.
  - ZOOKEEPER\_TICK\_TIME: Interval for Zookeeper heartbeats.
- **Ports:** Exposes port 2181 for Zookeeper client connections.

### Kafka

- **Image:** confluentinc/cp-kafka:7.0.0
- **Dependencies:** Depends on the zookeeper container.
- **Environment Variables:**
  - KAFKA\_BROKER\_ID: Unique ID for the Kafka broker.
  - KAFKA\_ZOOKEEPER\_CONNECT: Connection string for Zookeeper.
  - KAFKA\_ADVERTISED\_LISTENERS: Advertised listener for Kafka to allow other services to connect.

- `KAFKA_OFFSETS_TOPIC_REPLICATION_FACTOR`: Specifies replication factor for internal Kafka topics.
- **Ports**: Exposes port 9092 for Kafka connections.

## Postgres

- **Image**: postgres:13
- **Environment Variables**:
  - `POSTGRES_USER`: Username for the Postgres database.
  - `POSTGRES_PASSWORD`: Password for the Postgres database.
  - `POSTGRES_DB`: Name of the database created on startup.
- **Ports**: Exposes port 5432 for database connections.
- **Volumes**: Uses an `init.sql` file to initialize the database with sample data.
- **Logging Options**:
  - `max-size`: Sets maximum log size.
  - `max-file`: Sets maximum number of log files to retain.

## Debezium

- **Image**: debezium/connect:1.8
- **Dependencies**: Depends on both kafka and postgres services.
- **Environment Variables**:
  - `BOOTSTRAP_SERVERS`: Kafka broker address.
  - `GROUP_ID`: Consumer group ID for Debezium.
  - `CONFIG_STORAGE_TOPIC`, `OFFSET_STORAGE_TOPIC`, `STATUS_STORAGE_TOPIC`: Kafka topics for storing Debezium configuration, offsets, and status.
- **Ports**: Exposes port 8083 for Debezium REST API.

## Starting the Services

To start the services defined in the Docker Compose file, run:

```
docker-compose up -d
```

This command runs the containers in detached mode.

## Registering a Debezium Connector

Once all services are up and running, register a Debezium connector to monitor changes in Postgres. The following example assumes Debezium listens to the `testdb` database.

```
curl -i -X POST -H "Accept:application/json" -H "Content-Type:application/json" \
```

```

http://localhost:8083/connectors/\
-d '{
  "name": "postgres-connector",
  "config": {
    "connector.class":
"io.debezium.connector.postgresql.PostgresConnector",
    "tasks.max": "1",
    "database.hostname": "postgres",
    "database.port": "5432",
    "database.user": "postgres",
    "database.password": "password",
    "database.dbname": "testdb",
    "database.server.name": "pg-changes",
    "table.include.list": "public.api",
    "plugin.name": "pgoutput"
  }
}'

```

## Explanation of the Connector Configuration

- **connector.class:** Specifies the Debezium connector for PostgreSQL.
- **tasks.max:** Number of tasks to run concurrently.
- **database.hostname:** Hostname of the Postgres container.
- **database.port:** Port for Postgres.
- **database.user** and **database.password:** Credentials for connecting to Postgres.
- **database.dbname:** Name of the monitored database.
- **database.server.name:** Logical name used for topic naming in Kafka.
- **table.include.list:** Specifies which tables to monitor.
- **plugin.name:** Sets the output plugin, pgoutput for PostgreSQL.

## Verifying Data Capture in Kafka

After configuring Debezium, changes in the public.api table in Postgres will be streamed to the Kafka topic pg-changes.public.api.

You can use kafka-console-consumer to check messages in the topic:

```

docker exec -it kafka kafka-console-consumer --bootstrap-server
kafka:9092 --topic pg-changes.public.api --from-beginning

```

## Testing the POC

1. **Insert, Update, or Delete** data in the public.api/public.api\_usage table in Postgres using api's from baas-analytical-module.
2. **Monitor** the Kafka topic for change events.
3. **Verify** the CDC pipeline by observing the streamed data in Kafka.
4. Observe the message consumed by baas-charges-module.

## Stopping and Cleaning Up

To stop and remove the containers, use:

```
docker-compose down
```